

Using MOM6

Alistair Adcroft

GFDL MOM6 team includes Raphael Dussin, Robert Hallberg, Stephen Griffies, Matthew Harrison, Hae-Cheol Kim, John Krasting, Marshall Ward, Niki Zadeh

Other contributors for presentation: Kate Hedstrom, Andrew Shao



- “Using MOM6”
 - “Where to start?”
 - Every application is unique
 - Sometimes requires custom code specific to that one configuration
 - Building new configurations
 - Best to start from working example
1. Getting code/compiling/running
 2. Controlling
 - Parameters
 - Diagnostics
 3. Examples
 - Places to start
 4. Model/repository structure

“Getting started” (Cloning, compiling, running)

- “MOM6 wiki”
 - <https://github.com/NOAA-GFDL/MOM6-examples/wiki>
 - (first result on google)
- Instructions for
 - Cloning (obtaining code)
 - Compiling (fairly portable)
 - Running (fairly standard)
- User-contributed
 - Instructions assume some familiarity with linux & models
- Best limited to stand-alone ocean-only configurations
 - Coupled models are so much more complicated
- Does not cover working within an environment like GFDL’s FRE or CESM’s CIME
 - just low-level basics

Model input: run-time parameters

- Parameter syntax is key, value pairs

`KH = 25.`

- Self-documenting runs

```
815 KH = 25.0                ! [m2 s-1] default = 0.0
816                          ! The background Laplacian horizontal viscosity.
```

- Simple API

```
1429 call get_param(param_file, mdl, "KH", Kh,          &
1430               "The background Laplacian horizontal viscosity.", &
1431               units = "m2 s-1", default=0.0, scale=US%m_to_L**2*US%T_to_s)
```

- MOM6 always writes out

- `MOM_parameter_doc.all` (everything)
- `MOM_parameter_doc.short` (non-defaults)

- Bootstrapped parameter parser
 - namelist in `input.nml`

```
1  &MOM_input_nml
2      output_directory = '.',
3      input_filename = 'n'
4      restart_input_dir = 'INPUT',
5      restart_output_dir = 'RESTART',
6      parameter_filename = 'MOM_input',
7                          'MOM_salrestore',
8                          'MOM_override'
9  /
```

- Typical setup

- Baseline uses blank `MOM_override`
- Perturbation runs concisely contained in `MOM_override`

- Lots of error checking

Parallel decomposition

- Reported in MOM_parameter_doc.layout

```
22 NIHALO = 4      ! default = 4
23                ! The number of halo points on each side in the x-direction. With
24                ! STATIC_MEMORY_ this is set as NIHALO_ in MOM_memory.h at compile time; without
25                ! STATIC_MEMORY_ the default is NIHALO_ in MOM_memory.h (if defined) or 2.
26 NJHALO = 4      ! default = 4
27                ! The number of halo points on each side in the y-direction. With
28                ! STATIC_MEMORY_ this is set as NJHALO_ in MOM_memory.h at compile time; without
29                ! STATIC_MEMORY_ the default is NJHALO_ in MOM_memory.h (if defined) or 2.
30 NIPROC = 2      !
31                ! The number of processors in the x-direction. With STATIC_MEMORY_ this is set
32                ! in MOM_memory.h at compile time.
33 NJPROC = 4      !
34                ! The number of processors in the y-direction. With STATIC_MEMORY_ this is set
35                ! in MOM_memory.h at compile time.
36 LAYOUT = 2, 4   !
37                ! The processor layout that was actually used.
38 IO_LAYOUT = 1, 1 ! default = 1
39                ! The processor layout to be used, or 0,0 to automatically set the io_layout to
40                ! be the same as the layout.
```

- Bitwise reproduces across layout
- Optimal tile size $\sim 12 \times 12$ - 30×30
 - Tile size =
$$\text{NIGLOBAL}/\text{NIPROC}, \text{NJGLOBAL}/\text{NJPROC}$$
- Halos ~ 3 - 4 most typically needed
 - Parameters NIHALO, NJHALO
 - Without high-order advection and certain choices of time-stepping
- Tile dimensions should be \geq NIHALO, NJHALO
- Tiles may not be uniform!
 - Different sized tiles are allowed

Restarts

- Many runs take longer than a single job submission
- Bitwise reproducibility across a restart boundary
- Online time-averaged diagnostics are not included in restarts
 - handled by FMS framework
- Input data usually read from INPUT/
- Diagnostic output is in current directory
- Restart files generally written to RESTART/
- Restart files are read from INPUT/

Controlling diagnostics

- FMS diag_manager parses `diag_table`

- File definition

```
3  "ocean_daily",      1, "days",  1, "days", "time"
4  "ocean_month_snap", 1, "months", 1, "days", "time"
5  "ocean_month",      1, "months", 1, "days", "time"
6  "ocean_month_z",    1, "months", 1, "days", "time"
7  "ocean_annual",     12, "months", 1, "days", "time"
8  "ocean_annual_z",   12, "months", 1, "days", "time"
9  "ocean_scalar_month", 1, "months", 1, "days", "time"
10 "ocean_scalar_annual", 12, "months", 1, "days", "time"
11 "ocean_static",     -1, "months", 1, "days", "time"
```

- Variable lists per file/module

```
85  "ocean_model",  "thetao",      "thetao",      "ocean_annual",      "all", "mean", "none",2 # if use pre-TEOS10
86  #"ocean_model", "thetao",      "thetao",      "ocean_month",       "all", "mean", "none",2 # if use pre-TEOS10
87  "ocean_model_z", "thetao",      "thetao",      "ocean_annual_z",    "all", "mean", "none",2 # if use pre-TEOS10
88  "ocean_model_z", "thetao",      "thetao",      "ocean_month_z",     "all", "mean", "none",2 # if use pre-TEOS10
89  "ocean_model_z", "thetao_xyave", "thetao_xyave", "ocean_annual_z",    "all", "mean", "none",2 # if use pre-TEOS10
```

- [MOM6](#) wraps diag_manager
 - [Registers](#) same diagnostic in multiple vertical coordinates, multiple names, xy-averages
- Available diagnostics written by MOM6 at run-time
[available diags.0000](#)
- [Regional](#) diagnostics

MOM6-examples

```
MOM6-examples
├── coupled_AM2_LM3_SIS
├── coupled_AM2_LM3_SIS2
├── ice_ocean_SIS
├── ice_ocean_SIS2
│   ├── Baltic_OM4_025
│   ├── Baltic_OM4_05
│   ├── OM4_025
│   ├── OM4_05
│   └── SIS2_icebergs
├── ocean_only
│   ├── CVMix_SCM_tests
│   ├── DOME
│   ├── ISOMIP
│   ├── Phillips_2layer
│   ├── SCM_idealized_hurricane
│   ├── adjustment2d
│   ├── benchmark
│   ├── buoy_forced_basin
│   ├── circle_obcs
│   ├── double_gyre
│   ├── external_gwave
│   ├── flow_downslope
│   ├── global
│   ├── global_ALE
│   ├── idealized_hurricane
│   ├── lock_exchange
│   ├── mixed_layer_restrat_2d
│   ├── nonBous_global
│   ├── resting
│   ├── rotating_gravity_current
│   ├── seamount
│   ├── single_column
│   ├── sloshing
│   ├── tides_025
│   ├── torus_advection_test
│   ├── tracer_mixing
│   └── unit_tests
├── src
│   ├── FMS
│   ├── MOM6
│   ├── SIS2
│   ├── atmos_null
│   ├── coupler
│   ├── ice_param
│   ├── icebergs
│   ├── land_null
│   └── mkmf
├── tools
├── analysis
├── matlab
├── python
└── tests
```

ice-ocean configurations

ocean only configurations

sub-modules to source

- **double_gyre**
 - Wind driven gyre using stacked shallow water equations
- **Phillips_2layer**
 - Idealized channel model
- **flow_downslope**
 - adjustment problem over topography using different coordinates
- **OM_05**
 - 1/2° ice-ocean global model

ocean_only/double_gyre

Visualizing and animating sea-surface height in the double-gyre example.ipynb

We will use matplotlib.pyplot for plotting and scipy's netcdf package for reading the model output. The %pylab inline causes figures to appear in the page and conveniently alias pyplot to plt (which is becoming a widely used alias).

This analysis assumes you changed DAYMAX to some multiple of 5 so that there are multiple time records in the model output.

To see this notebook with figures, see <https://gist.github.com/adcroft/2a2b91d66625fd534372>.

```
In [1]: %pylab inline
import scipy.io.netcdf

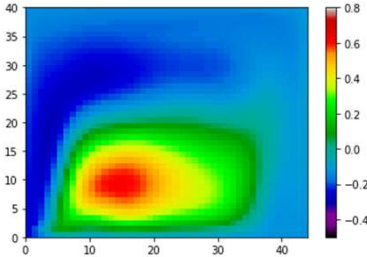
Populating the interactive namespace from numpy and matplotlib

We first create a netcdf object, or "handle", to the netcdf file. We'll also list all the objects in the netcdf object.
```

```
In [2]: prog_file = scipy.io.netcdf_file('prog_0001_006.nc')
prog_file.variables
```

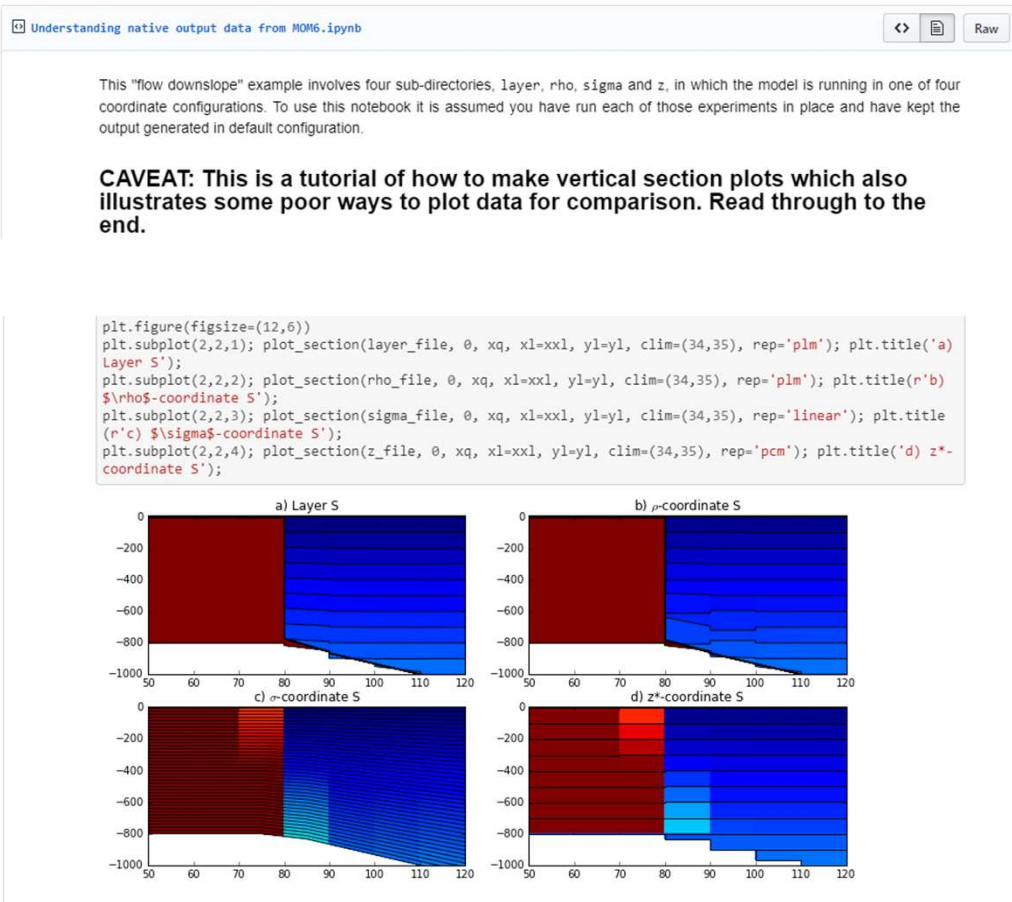
```
Out[2]: {'Time': <scipy.io.netcdf.netcdf_variable at 0xb60364ac>,
'e': <scipy.io.netcdf.netcdf_variable at 0xb603686c>,
'h': <scipy.io.netcdf.netcdf_variable at 0xb603676c>,
'u': <scipy.io.netcdf.netcdf_variable at 0xb603668c>,
'v': <scipy.io.netcdf.netcdf_variable at 0xb603670c>,
'xh': <scipy.io.netcdf.netcdf_variable at 0xb603648c>,
'xq': <scipy.io.netcdf.netcdf_variable at 0xb603634c>,
'yh': <scipy.io.netcdf.netcdf_variable at 0xb603632c>,
'yq': <scipy.io.netcdf.netcdf_variable at 0xb603652c>,
'zi': <scipy.io.netcdf.netcdf_variable at 0xb60365ec>,
'zl': <scipy.io.netcdf.netcdf_variable at 0xb60363cc>}
```

```
In [10]: for n in range( e_handle.shape[0]):
display.display(plt.gcf())
plt.clf()
plot_ssh(n)
display.clear_output(wait=True)
```



- jupyter notebook
 - (sorry about rainbow colormap)
 - This one uses matplotlib and scipy
- Much more needs to be added
 - Notebook for other examples
 - Others use netCDF4 instead of scipy
 - Will add xarray, seaborn examples
- Very much NOT advocating for one analysis system/style

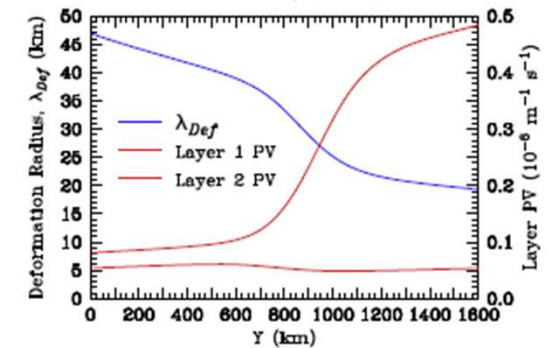
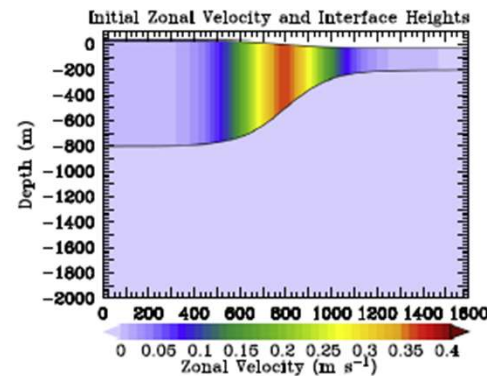
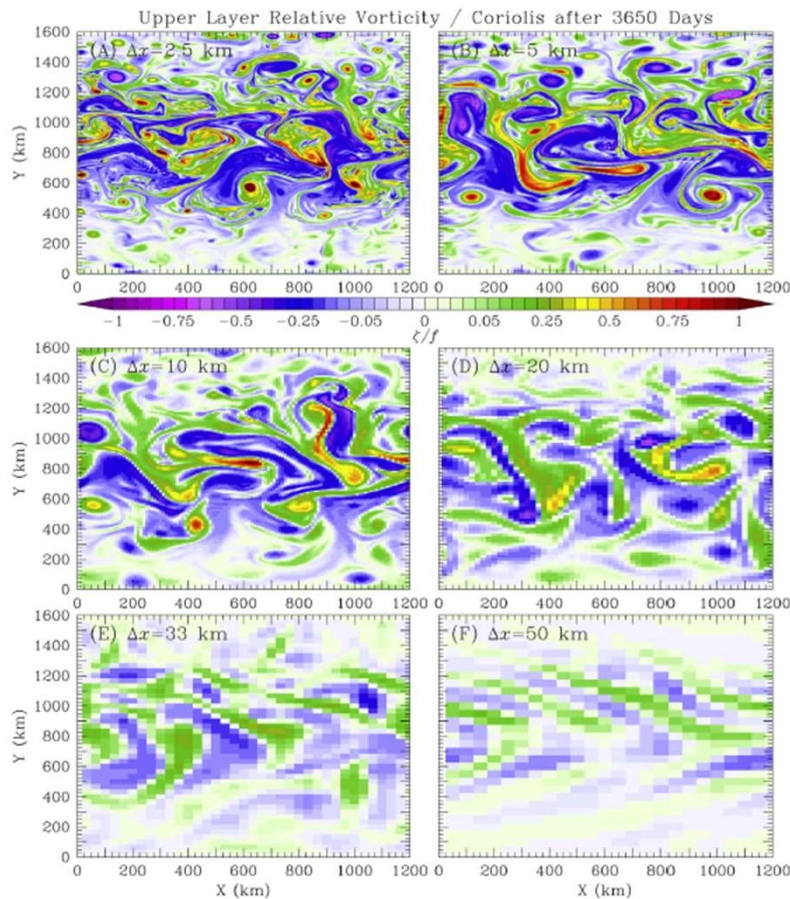
ocean_only/flow_downslope



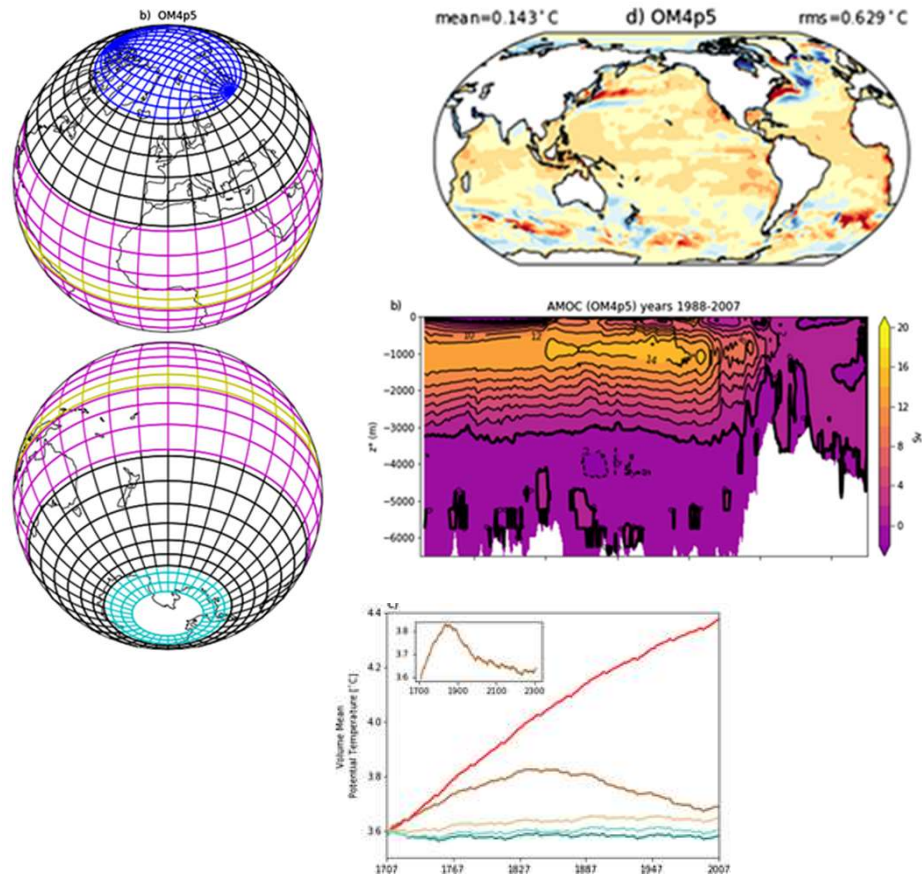
- 2d density current
- Notebook is a treatise on how plotting in the vertical can go wrong
 - Idea is to explain how the model stores data in the vertical
 - ... and how to look at the vertical in native space

ocean_only/Phillips_2layer

- Hallberg, 2013
- Idealized zonal channel
 - Customized forcing
- No jupyter notebooks yet
 - plots were done with Ferret



ice_ocean_SIS2/OM4_05


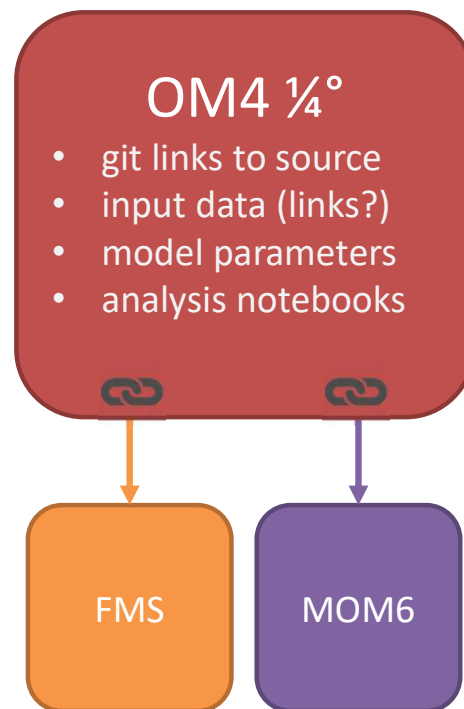


- $\frac{1}{2}^\circ$ global ice-ocean model
 - Uses GFDL SIS2 sea-ice model and GFDL coupler
- Uses GFDL vertical physics
 - ePBL, JHL, ...
- Non-eddying (coarse resolution)
- Uses GM and neutral-diffusion parameterizations

Adcroft et al., 2019

Experiment oriented repositories

- Using a repository for experiment development
 - treating configurations like code
- No new tools (just `git`)
- Provides history of experiment design
 - Recoverable / reproducible
- Used in other workflows
 - e.g. Payu, ROMS



Repository organization (experiment suites)

- Layered repositories using sub-modules

- Regression results

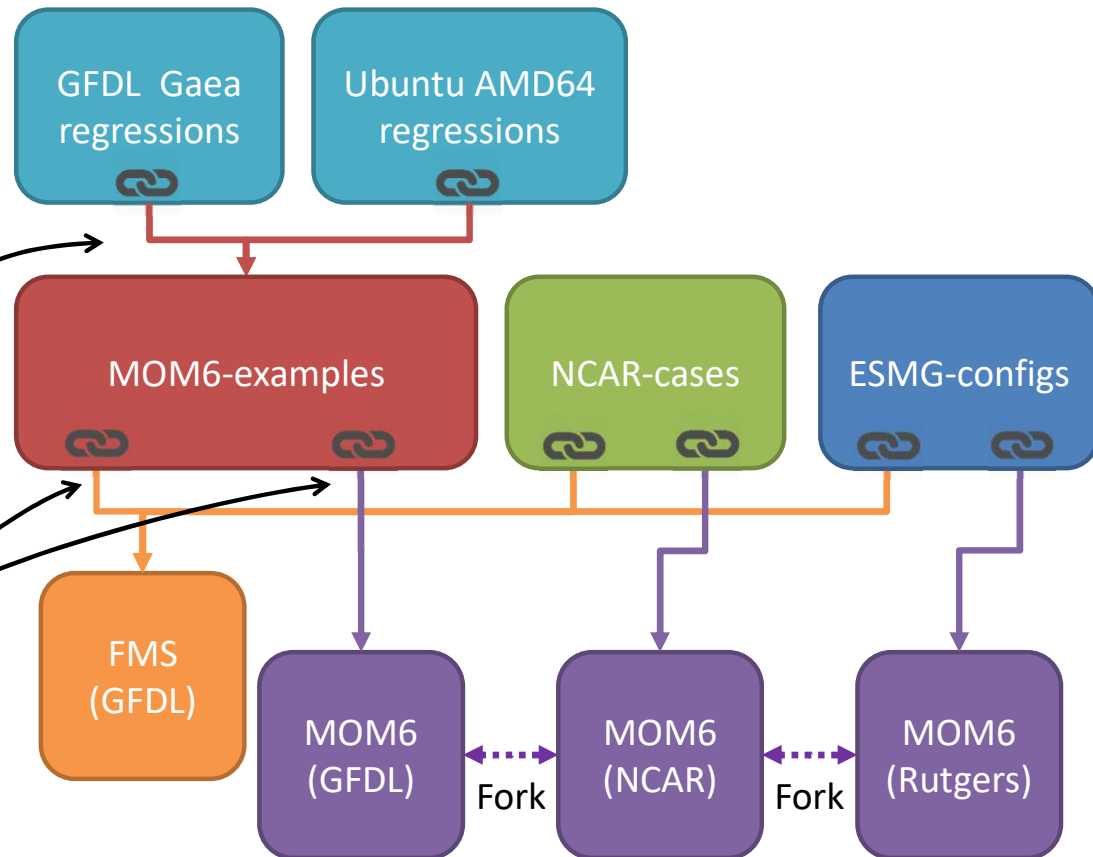
- Output from regression tests
 - Platform dependent
- Records specific version of configurations

- Configurations

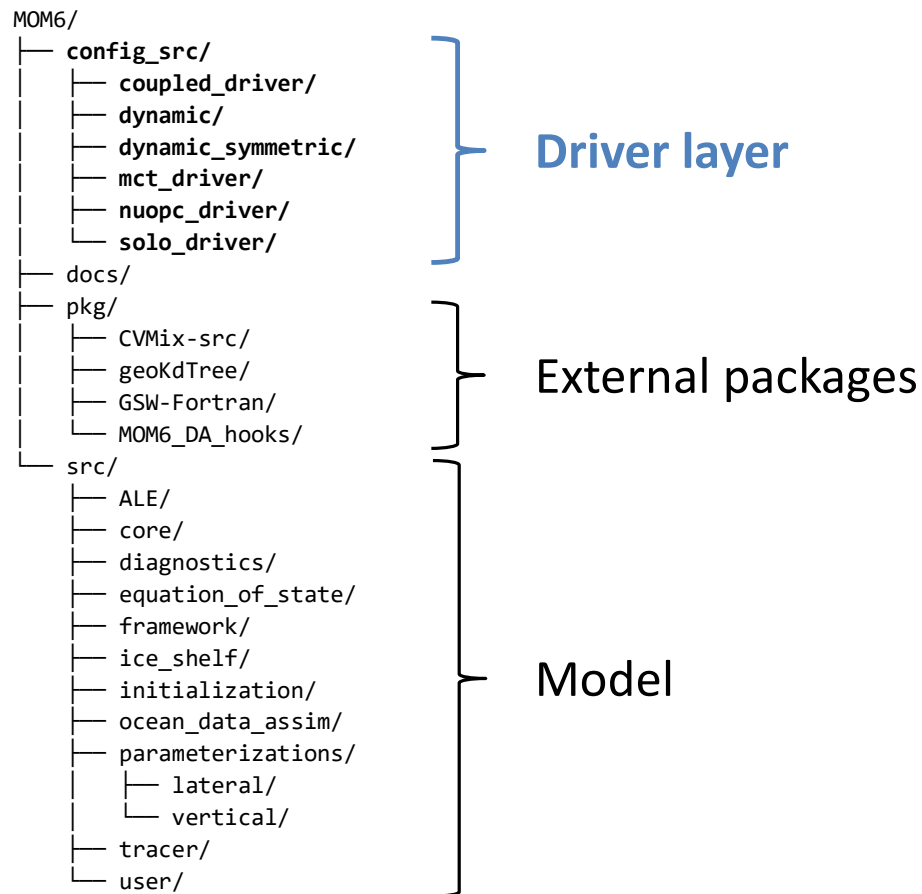
- Input files (parameters)
- Records specific versions of source
 - Including URLs (for forks)

- Source for MOM6, FMS, SIS2, ...

- Pure source code (+ packages)



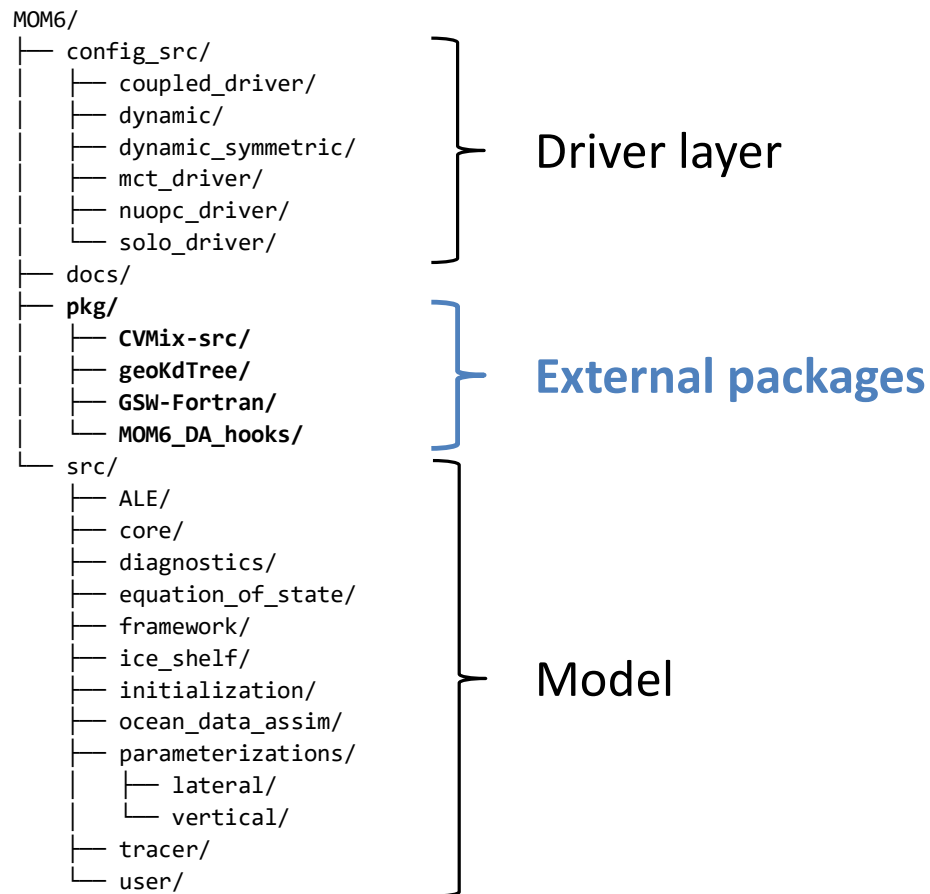
Code tree: driver layer



config_src/

- Selectively compiled
 - NCAR coupled mode:
nuopc_driver + dynamic
 - Stand-alone ocean model
solo_driver + dynamic
- Alternative version of same code
e.g.
dynamic or dynamic_symmetric

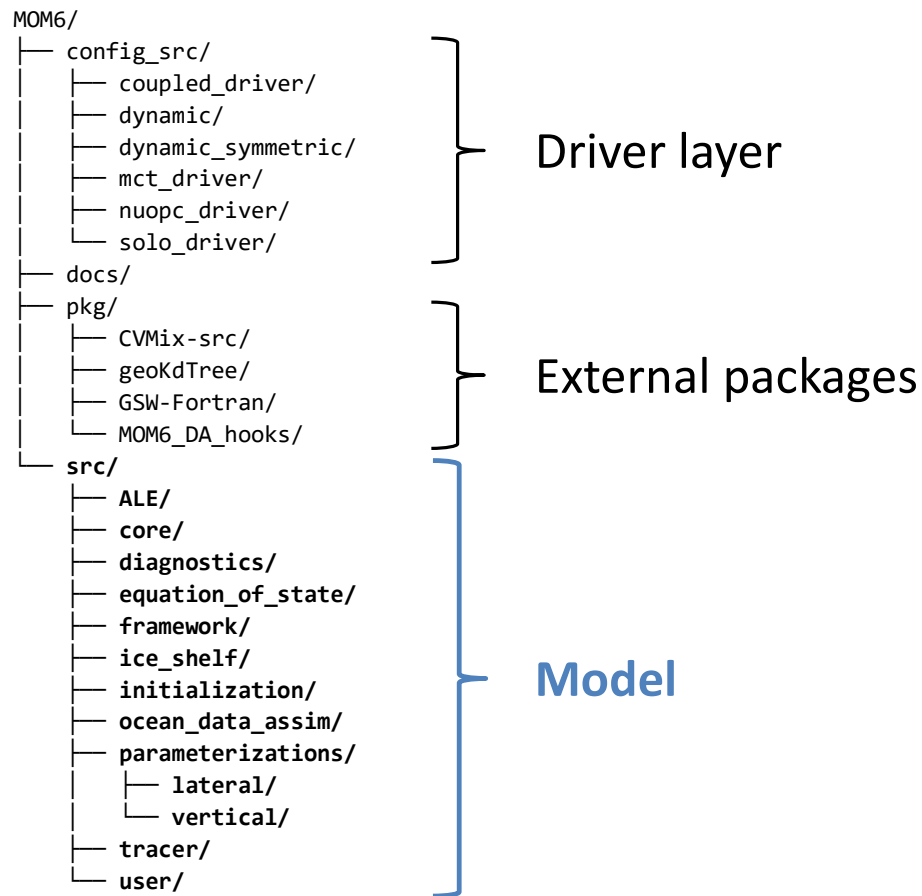
Code tree: external packages



pkg/

- Not compiled in place
- Symbolic links to required source live under `src/` and point to `pkg/`
 - External packages often contain more than source and not all source compiles!
- Each package is a `git` submodule
 - using specific commit hash

Code tree: main model



src/

- Code for solving equations of motion, tracers, diagnostics, etc.
- Always compiled
- No CPP macros except for MEMORY and GRID macros
 - the few existing exceptions will be removed one day

Code tree: main model

MOM6/

└─ src/

├─ ALE/

- Vertical remapping

├─ core/

- Dynamic core (momentum, continuity)

├─ diagnostics/

- Some collective diagnostic

├─ equation_of_state/

- Five equations of state

├─ framework/

- Interface to FMS (communications, I/O)

├─ initialization/

- Grid/state allocation/initialization

├─ parameterizations/

├─ └─ lateral/

- Lateral parameterizations

├─ └─ vertical/

- Vertical parameterizations

├─ tracer/

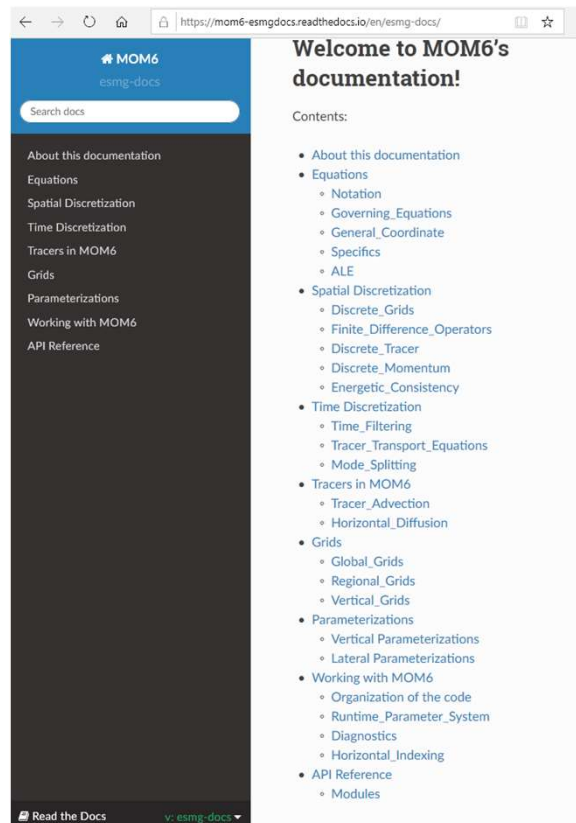
- Tracers including parameterizations (Redi)

└─ user/

- Configuration specific code (forcing/init)

Code tree: the most important bits

MOM6/
├── .testing/
└── docs/



.testing/

- Continuous integration
 - Runs tests on Travis-CI (soon also GitHub Actions)
 - Can be used for development

docs/

- Source for documentation hosted at <https://mom6.readthedocs.io>
 - Under dev. by K. Hedstrom

Future topics

- Verification and validation of MOM6 contributions
 - Marshall Ward
- Equations and algorithms
 - Bob Hallberg
- Lagrangian remap method
 - Stephen Griffies
- Analysis and tools
 - Raphael Dussin
- Ocean data assimilation interfaces
 - Matthew Harrison